


```
In [1]: import pandas as pd
import pickle
import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: a=pd.read_csv("C:\\Users\\reshma_koduri\\OneDrive\\Documents\\archive (2)\\voice.csv")
a
```

```
Out[2]:
```

	meanfreq	sd	median	Q25	Q75	IQR	skew	kurt	sp.ent
0	0.059781	0.064241	0.032027	0.015071	0.090193	0.075122	12.863462	274.402906	0.893369
1	0.066009	0.067310	0.040229	0.019414	0.092666	0.073252	22.423285	634.613855	0.892193
2	0.077316	0.083829	0.036718	0.008701	0.131908	0.123207	30.757155	1024.927705	0.846389
3	0.151228	0.072111	0.158011	0.096582	0.207955	0.111374	1.232831	4.177296	0.963322
4	0.135120	0.079146	0.124656	0.078720	0.206045	0.127325	1.101174	4.333713	0.971955
...
3163	0.131884	0.084734	0.153707	0.049285	0.201144	0.151859	1.762129	6.630383	0.962934
3164	0.116221	0.089221	0.076758	0.042718	0.204911	0.162193	0.693730	2.503954	0.960716
3165	0.142056	0.095798	0.183731	0.033424	0.224360	0.190936	1.876502	6.604509	0.946854
3166	0.143659	0.090628	0.184976	0.043508	0.219943	0.176435	1.591065	5.388298	0.950436
3167	0.165509	0.092884	0.183044	0.070072	0.250827	0.180756	1.705029	5.769115	0.938829

3168 rows × 21 columns



```
In [3]: a.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3168 entries, 0 to 3167
Data columns (total 21 columns):
#   Column      Non-Null Count  Dtype
---  -
0   meanfreq    3168 non-null   float64
1   sd          3168 non-null   float64
2   median      3168 non-null   float64
3   Q25         3168 non-null   float64
4   Q75         3168 non-null   float64
5   IQR         3168 non-null   float64
6   skew        3168 non-null   float64
7   kurt        3168 non-null   float64
8   sp.ent      3168 non-null   float64
9   sfm         3168 non-null   float64
10  mode        3168 non-null   float64
11  centroid    3168 non-null   float64
12  meanfun     3168 non-null   float64
13  minfun      3168 non-null   float64
14  maxfun      3168 non-null   float64
15  meandom     3168 non-null   float64
16  mindom      3168 non-null   float64
17  maxdom      3168 non-null   float64
18  dfrange     3168 non-null   float64
```

12/12/23, 11:13 AM

gender recognition by voice

19 modindx 3168 non-null float64

20 label 3168 non-null object

dtypes: float64(20), object(1)

memory usage: 519.9+ KB

In [4]:

a.describe()

Out[4]:

	meanfreq	sd	median	Q25	Q75	IQR	skew	
count	3168.000000	3168.000000	3168.000000	3168.000000	3168.000000	3168.000000	3168.000000	3
mean	0.180907	0.057126	0.185621	0.140456	0.224765	0.084309	3.140168	
std	0.029918	0.016652	0.036360	0.048680	0.023639	0.042783	4.240529	
min	0.039363	0.018363	0.010975	0.000229	0.042946	0.014558	0.141735	
25%	0.163662	0.041954	0.169593	0.111087	0.208747	0.042560	1.649569	
50%	0.184838	0.059155	0.190032	0.140286	0.225684	0.094280	2.197101	
75%	0.199146	0.067020	0.210618	0.175939	0.243660	0.114175	2.931694	
max	0.251124	0.115273	0.261224	0.247347	0.273469	0.252225	34.725453	1

In [5]:

list(a)

Out[5]:

['meanfreq',
'sd',
'median',
'Q25',
'Q75',
'IQR',
'skew',
'kurt',
'sp.ent',
'sfm',
'mode',
'centroid',
'meanfun',
'minfun',
'maxfun',
'meandom',
'mindom',
'maxdom',
'dfrange',
'modindx',
'label']

In [6]:

a.isna().sum()

Out[6]:

meanfreq 0
sd 0
median 0
Q25 0
Q75 0
IQR 0
skew 0
kurt 0
sp.ent 0
sfm 0

```

mode          0
centroid      0
meanfun       0
minfun        0
maxfun        0
meandom       0
mindom        0
maxdom        0
dfrange       0
modindx       0
label         0
dtype: int64

```

In [7]:

```

b=a.drop(['Q25', 'Q75', 'IQR', 'kurt', 'sp.ent', 'sfm', 'centroid', 'skew', 'sd', 'modindx'],
b

```

Out[7]:

	meanfreq	median	mode	meanfun	minfun	maxfun	meandom	mindom	maxdom	d
0	0.059781	0.032027	0.000000	0.084279	0.015702	0.275862	0.007812	0.007812	0.007812	0.
1	0.066009	0.040229	0.000000	0.107937	0.015826	0.250000	0.009014	0.007812	0.054688	0.
2	0.077316	0.036718	0.000000	0.098706	0.015656	0.271186	0.007990	0.007812	0.015625	0.
3	0.151228	0.158011	0.083878	0.088965	0.017798	0.250000	0.201497	0.007812	0.562500	0.
4	0.135120	0.124656	0.104261	0.106398	0.016931	0.266667	0.712812	0.007812	5.484375	5.
...
3163	0.131884	0.153707	0.200836	0.182790	0.083770	0.262295	0.832899	0.007812	4.210938	4.
3164	0.116221	0.076758	0.013683	0.188980	0.034409	0.275862	0.909856	0.039062	3.679688	3.
3165	0.142056	0.183731	0.008006	0.209918	0.039506	0.275862	0.494271	0.007812	2.937500	2.
3166	0.143659	0.184976	0.212202	0.172375	0.034483	0.250000	0.791360	0.007812	3.593750	3.
3167	0.165509	0.183044	0.267702	0.185607	0.062257	0.271186	0.227022	0.007812	0.554688	0.

3168 rows × 11 columns



In [8]:

```

c=pd.get_dummies(b,dtype=int)
c

```

Out[8]:

	meanfreq	median	mode	meanfun	minfun	maxfun	meandom	mindom	maxdom	d
0	0.059781	0.032027	0.000000	0.084279	0.015702	0.275862	0.007812	0.007812	0.007812	0.
1	0.066009	0.040229	0.000000	0.107937	0.015826	0.250000	0.009014	0.007812	0.054688	0.
2	0.077316	0.036718	0.000000	0.098706	0.015656	0.271186	0.007990	0.007812	0.015625	0.
3	0.151228	0.158011	0.083878	0.088965	0.017798	0.250000	0.201497	0.007812	0.562500	0.
4	0.135120	0.124656	0.104261	0.106398	0.016931	0.266667	0.712812	0.007812	5.484375	5.
...
3163	0.131884	0.153707	0.200836	0.182790	0.083770	0.262295	0.832899	0.007812	4.210938	4.
3164	0.116221	0.076758	0.013683	0.188980	0.034409	0.275862	0.909856	0.039062	3.679688	3.
3165	0.142056	0.183731	0.008006	0.209918	0.039506	0.275862	0.494271	0.007812	2.937500	2.

	meanfreq	median	mode	meanfun	minfun	maxfun	meandom	mindom	maxdom	d
3166	0.143659	0.184976	0.212202	0.172375	0.034483	0.250000	0.791360	0.007812	3.593750	3.
3167	0.165509	0.183044	0.267702	0.185607	0.062257	0.271186	0.227022	0.007812	0.554688	0.

3168 rows × 12 columns

```
In [9]: y=c['meanfreq']
y
```

```
Out[9]: 0      0.059781
1      0.066009
2      0.077316
3      0.151228
4      0.135120
...
3163   0.131884
3164   0.116221
3165   0.142056
3166   0.143659
3167   0.165509
Name: meanfreq, Length: 3168, dtype: float64
```

```
In [10]: x=c.drop(['meanfreq'],axis=1)
x
```

```
Out[10]:
```

	median	mode	meanfun	minfun	maxfun	meandom	mindom	maxdom	dfrange	lab
0	0.032027	0.000000	0.084279	0.015702	0.275862	0.007812	0.007812	0.007812	0.000000	
1	0.040229	0.000000	0.107937	0.015826	0.250000	0.009014	0.007812	0.054688	0.046875	
2	0.036718	0.000000	0.098706	0.015656	0.271186	0.007990	0.007812	0.015625	0.007812	
3	0.158011	0.083878	0.088965	0.017798	0.250000	0.201497	0.007812	0.562500	0.554688	
4	0.124656	0.104261	0.106398	0.016931	0.266667	0.712812	0.007812	5.484375	5.476562	
...
3163	0.153707	0.200836	0.182790	0.083770	0.262295	0.832899	0.007812	4.210938	4.203125	
3164	0.076758	0.013683	0.188980	0.034409	0.275862	0.909856	0.039062	3.679688	3.640625	
3165	0.183731	0.008006	0.209918	0.039506	0.275862	0.494271	0.007812	2.937500	2.929688	
3166	0.184976	0.212202	0.172375	0.034483	0.250000	0.791360	0.007812	3.593750	3.585938	
3167	0.183044	0.267702	0.185607	0.062257	0.271186	0.227022	0.007812	0.554688	0.546875	

3168 rows × 11 columns



```
In [11]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33,random_state=42)
```

```
In [12]: from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestRegressor
reg=RandomForestRegressor()
```

```
n_estimators=[25,50,75,100,125,150,175,200]
criterion=['mse']
max_depth=[3,5,10]
parameters={'n_estimators': n_estimators, 'criterion': criterion, 'max_depth': max_depth}
rfc_reg = GridSearchCV(reg, parameters)
rfc_reg.fit(x_train,y_train)
```

```
Out[12]: GridSearchCV(estimator=RandomForestRegressor(),
                  param_grid={'criterion': ['mse'], 'max_depth': [3, 5, 10],
                              'n_estimators': [25, 50, 75, 100, 125, 150, 175, 200]})
```

```
In [13]: rfc_reg.best_params_
```

```
Out[13]: {'criterion': 'mse', 'max_depth': 10, 'n_estimators': 150}
```

```
In [23]: reg=RandomForestRegressor(n_estimators=150,criterion='mse',max_depth=10)
```

```
In [24]: reg.fit(x_train,y_train)
```

```
Out[24]: RandomForestRegressor(max_depth=10, n_estimators=150)
```

```
In [31]: y_pred=reg.predict(x_test)
y_pred
```

```
Out[31]: array([0.17568116, 0.1906409 , 0.15965559, ..., 0.22959718, 0.21601968,
                0.18506223])
```

```
In [32]: from sklearn.metrics import r2_score
r2_score(y_test,y_pred)
```

```
Out[32]: 0.9477014545373129
```

```
In [37]: Results= pd.DataFrame(columns=['actual','Predicted'])
Results['actual']=y_test
Results['Predicted']=y_pred
Results=Results.reset_index()
Results['Id']=Results.index
Results.head(10)
```


```
Out[37]:
```

	index	actual	Predicted	Id
0	2148	0.186833	0.175681	0
1	1124	0.188879	0.190641	1
2	170	0.150705	0.159656	2
3	3158	0.183667	0.172573	3
4	2229	0.205159	0.201107	4
5	1960	0.141798	0.149672	5
6	411	0.186064	0.190510	6
7	457	0.138354	0.138139	7
8	2881	0.201499	0.196011	8

	index	actual	Predicted	Id
9	602	0.169100	0.169062	9

```
In [39]: import seaborn as sns
import matplotlib.pyplot as plt
sns.lineplot(x='Id',y='actual',data=Results.head(50))
sns.lineplot(x='Id',y='Predicted',data=Results.head(50))
plt.plot()
```

Out[39]: []



In []: